

Winning Space Race with Data Science

Applied Data Science Capstone

Oscar, Kuan Ka Meng
Jan 18, 2022



Outline

- Executive Summary
- Introduction
- Methodology
- Results
- Conclusion
- Appendix

Executive Summary

- Summary of methodologies
 - Data Collection
 - Data Wrangling
 - EDA with Data Visualization
 - EDA with SQL
 - Building an Interactive Map with Folium
 - Building a Dashboard with Plotly Dash
 - Predictive Analysis (Classification)
- Summary of all results
 - Exploratory Data Analysis Results
 - Interactive Analytics Demo in Screenshots
 - Predictive Analysis Results

Introduction

- Project background and context

SpaceX is the company that are making space-travelling for everyone. Their rocket launch is relatively common and not expensive.

Also, *SpaceX* shows *Falcon 9* rocket launches on its website, with a cost of 62 million dollars; other providers cost upward of 165 million dollars each.

Therefore, we predict the first stage of *Falcon 9* will land successfully. If our prediction is correct, we can determine the cost of a launch. This information can be used if an alternate company wants to bid against *SpaceX* for a rocket launch.

- Problems you want to find answers

- Correlations between each rocket variables and successful landing rate
- Conditions to get the best results and ensure the best successful landing rate

Section 1

Methodology

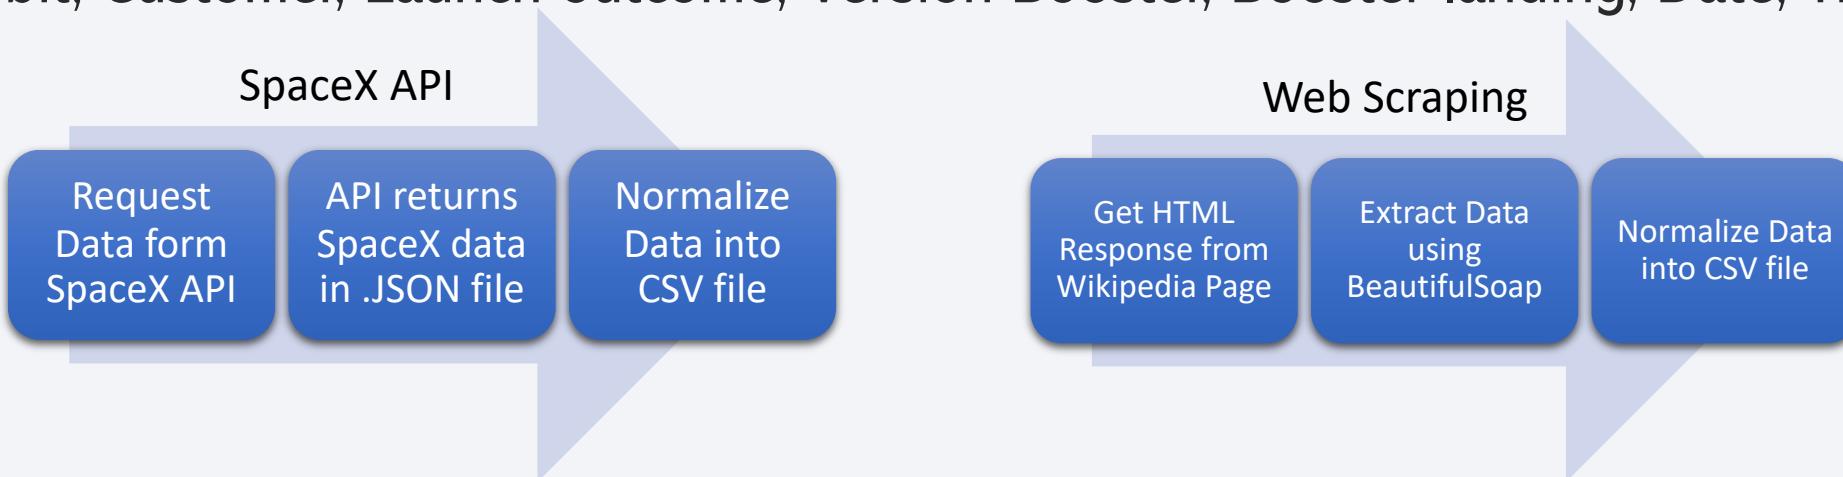
Methodology

Executive Summary

- Data collection methodology:
 - SpaceX API & Web Scraping – Falcon 9 and Falcon Heavy Launches Records from Wikipedia.
- Perform data wrangling
 - Convert outcomes into Training Labels with the booster successfully/unsuccessful landed
- Perform exploratory data analysis (EDA) using visualization and SQL
- Perform interactive visual analytics using Folium and Plotly Dash
- Perform predictive analysis using classification models
 - Find best Hyperparameter for SVM, Classification Trees and Logistic Regression

Data Collection

- The data collection process includes a combination of API requests from the SpaceX API and web scraping data from a table in the Wikipedia page of Space X, Falcon 9 and Falcon Heavy Launches Records.
- SpaceX API Data Columns: FlightNumber, Date, BoosterVersion, PayloadMass, Orbit, LaunchSite, Outcome, Flights, GridFins, Reused, Legs, LandingPad, Block, ReusedCount, Serial, Longitude, Latitude
- Wikipedia Web Scrape Data Columns: Flight No., Launch site, Payload, PayloadMass, Orbit, Customer, Launch outcome, Version Booster, Booster landing, Date, Time



Data Collection – SpaceX API

1. Requesting Rocket Launch Data from SpaceX API

```
spacex_url="https://api.spacexdata.com/v4/launches/past"

response = requests.get(spacex_url)
```

2. Converting Response to a JSON file

```
# Use json_normalize meethod to convert the json result into a dataframe
data = pd.json_normalize(response.json())
```

3. Using Custom Functions to Clean Data

```
# Call getBoosterVersion
getBoosterVersion(data)
```

```
# Call getLaunchSite
getLaunchSite(data)
```

```
# Call getPayloadData
getPayloadData(data)
```

```
# Call getCoreData
getCoreData(data)
```

GitHub Link:

https://github.com/sunshinekuan/IBM-Data-Science-Professional-Certificate/blob/main/Applied%20Data%20Science%20Capstone/01_Data%20Collection%20API.ipynb

4. Combining the Columns into a Dictionary to Create Data Frame

```
launch_dict = {'FlightNumber': list(data['flight_number']),
'Date': list(data['date']),
'BoosterVersion':BoosterVersion,
'PayloadMass':PayloadMass,
'Orbit':Orbit,
'LaunchSite':LaunchSite,
'Outcome':Outcome,
'Flights':Flights,
'GridFins':GridFins,
'Reused':Reused,
'Legs':Legs,
'LandingPad':LandingPad,
'Block':Block,
'ReusedCount':ReusedCount,
'Serial':Serial,
'Serial':Serial,
'Longitude': Longitude,
'Latitude': Latitude}
```

```
# Create a data from launch_dict

launch_df = pd.DataFrame.from_dict(launch_dict)
```

5. Filtering Data Frame and Exporting to a CSV file

```
# Hint data['BoosterVersion']!='Falcon 1'
data_falcon9 = launch_df[launch_df['BoosterVersion'] == 'Falcon 9']

data_falcon9.to_csv('dataset_part\1.csv', index=False)
```

Data Collection - Scraping

1. Getting Response from HTML

```
html_data = requests.get(static_url).text
```

2. Creating a BeautifulSoup Object

```
soup = BeautifulSoup(html_data, "html5lib")
```

3. Finding All Tables and Assigning the Result to a List

```
html_tables = soup.find_all('table')
```

4. Extracting Column Name One-By-One

```
column_names = []
for row in first_launch_table.find_all('th'):
    name = extract_column_from_header(row)
    if (name != None and len(name) > 0):
        column_names.append(name)
```

GitHub Link:

https://github.com/sunshinekuan/IBM-Data-Science-Professional-Certificate/blob/main/Applied%20Data%20Science%20Capstone/02_Data%20Collection%20with%20Web%20Scraping.ipynb

5. Creating an Empty Dictionary with Keys

```
launch_dict= dict.fromkeys(column_names)
del launch_dict['Date and time ( )']

# Let's initial the launch_dict with each value to be an empty list
launch_dict['Flight No.'] = []
launch_dict['Launch site'] = []
launch_dict['Payload'] = []
launch_dict['Payload mass'] = []
launch_dict['Orbit'] = []
launch_dict['Customer'] = []
launch_dict['Launch outcome'] = []

# Added some new columns
launch_dict['Version Booster']= []
launch_dict['Booster landing']= []
launch_dict['Date']= []
launch_dict['Time']= []
```

6. Filling up the *launch_dict* with launch records (Please Refer to the Notebook)

7. Creating a Dataframe and Exporting it to a CSV

```
df=pd.DataFrame(launch_dict)
df.to_csv('spacex_web_scraped.csv', index=False)
```

Data Wrangling

- There are several cases in which the booster failed to successfully land on the dataset. It attempted to land but failed because of accident sometimes.
 - *True Ocean*: the mission result is successfully landed in a specific area of the ocean
 - *False Ocean*: the mission result is not successfully landed in a specific area of the ocean
 - *True RTLS*: the mission result is successfully landed on the ground pad
 - *False RTLS*: the mission result is not successfully landed on the ground pad
 - *True ASDS*: the mission result is successfully landed on the drone ship
 - *False ASDS*: the mission result is successfully landed on the drone ship
- **Converting the above results into training labels:**
 - 1 = success
 - 0 = failure

Data Wrangling

1. Calculating the Number of Launches at Each Site

```
df[ 'LaunchSite' ].value_counts()
```

2. Calculating the Number and Occurrence of Each Orbit

```
df.Orbit.value_counts()
```

3. Calculating the Number and Occurrence of Mission Outcome per Orbit Type

```
landing_outcomes = df.Outcome.value_counts()
```

4. Creating a Landing Outcome Label from Outcome Column

```
landing_class = []
for outcome in df.Outcome:
    if outcome in bad_outcomes:
        landing_class.append(0)
    else:
        landing_class.append(1)

df[ 'Class' ]=landing_class
```

5. Calculating the Success Rate for Every Landing in Dataset

```
df[ "Class" ].mean()
```

```
0.6666666666666666
```

6. Exporting Dataset to a CSV

```
df.to_csv("dataset_part\2.csv", index=False)
```

GitHub Link:

https://github.com/sunshinekuan/IBM-Data-Science-Professional-Certificate/blob/main/Applied%20Data%20Science%20Capstone/03_EDA.ipynb

EDA with Data Visualization

- Scatter Chart
 - *Flight Number vs Launch Site*
 - *Payload vs Launch Site*
 - *Flight Number vs Orbit Type*
 - *Payload vs Orbit Type*

(A scatter chart shows how much one variable is affected by another. The relationship between two variables is called a correlation. This plot is generally composed of large data bodies.)

EDA with Data Visualization

- Bar Chart

- *Orbit Type vs Success Rate*

(A bar chart makes it easy to compare datasets between multiple groups. One axis represents a category, and the other axis represents a discrete value. The purpose of this chart is to indicate the relationship between the two axes.)

- Line Chart

- *Year vs Success Rate*

(A line chart shows data variables and trends very clearly and helps predict the results of data that has not yet been recorded.)

GitHub Link:

https://github.com/sunshinekuan/IBM-Data-Science-Professional-Certificate/blob/main/Applied%20Data%20Science%20Capstone/05_EDA%20with%20Visualization.ipynb

EDA with SQL

GitHub Link:

https://github.com/sunshinekuan/IBM-Data-Science-Professional-Certificate/blob/main/Applied%20Data%20Science%20Capstone/04_EDA%20with%20SQL.ipynb

- Loading the dataset into the corresponding table in a Db2 database, and executing SQL queries to answer the following questions
 1. Displaying the names of the unique launch sites in the space mission
 2. Displaying 5 records where launch sites begin with the string “CCA”
 3. Displaying the total payload mass carried by boosters launched by NASA (CRS)
 4. Displaying average payload mass carried by booster version F9 v1.1
 5. Listing the date when the first successful landing outcome in ground pad was achieved
 6. Listing the names of the boosters which is success in drone ship and have payload mass greater than 4000 but less than 6000
 7. Listing the total number of successful and failure mission outcomes
 8. Listing the names of *booster_versions* which have carried the maximum payload mass
 9. Listing the failed *landing_outcomes* in drone ship, their booster versions, and launch site names in year 2015
 10. Ranking the count of landing outcomes (such as Failure (drone ship) or Success (ground pad)) between the date 2010-06-04 and 2017-03-20, in descending order

Build an Interactive Map with Folium

- Objects created and added to a folium map
 - Markers that show all launch sites on a map
 - Markers that show the success/failed launches for each site on the map
 - Lines that show the distances between a launch site to its proximities
- By adding these objects, following geographical patterns about launch sites are found
 - Are launch sites in close proximity to railways? Yes
 - Are launch sites in close proximity to highways? Yes
 - Are launch sites in close proximity to coastline? Yes
 - Do launch sites keep certain distance away from cities? Yes

GitHub Link:

https://github.com/sunshinekuan/IBM-Data-Science-Professional-Certificate/blob/main/Applied%20Data%20Science%20Capstone/06_Interactive%20Visual%20Analytics%20with%20Folium.ipynb

Build a Dashboard with Plotly Dash

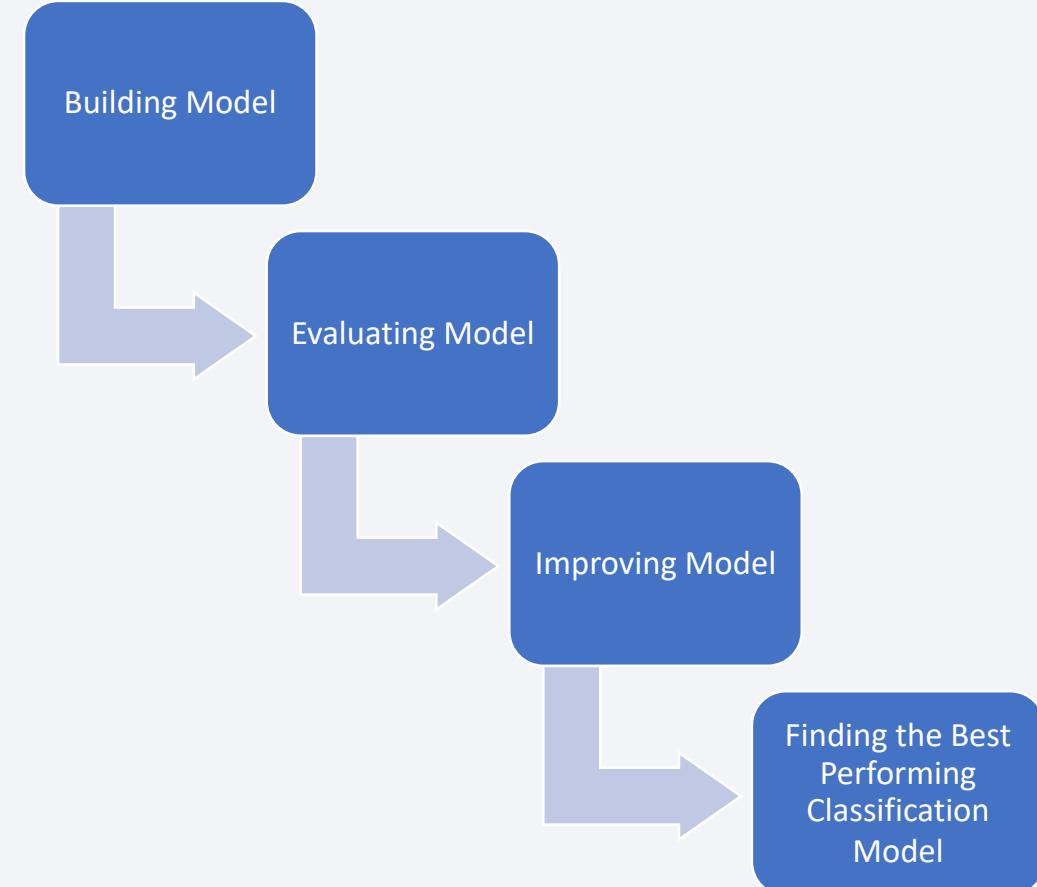
- The dashboard application contains a pie chart and a scatter point chart.
 - Pie chart
 - For showing total success launches by sites
 - This chart can be selected to indicate a successful landing distribution across all launch sites or to indicate the success rate of individual launch sites.
 - Scatter chart
 - For showing the relationship between Outcomes and Payload mass(Kg) by different boosters
 - Has 2 inputs: All sites/individual site & Payload mass on a slider between 0 and 10000 kg
 - This chart helps determine how success depends on the launch point, payload mass, and booster version categories.

GitHub Link (Code Only):

https://github.com/sunshinekuan/IBM-Data-Science-Professional-Certificate/blob/main/Applied%20Data%20Science%20Capstone/07_Interactive%20Dashboard%20with%20Ploty%20Dash.py

Predictive Analysis (Classification)

- Perform exploratory Data Analysis and determine Training Labels
 - Create a column for the class
 - Standardize the data
 - Split into training data and test data
- Find best Hyperparameter for SVM, Classification Trees and Logistic Regression
 - Find the method performs best using test data

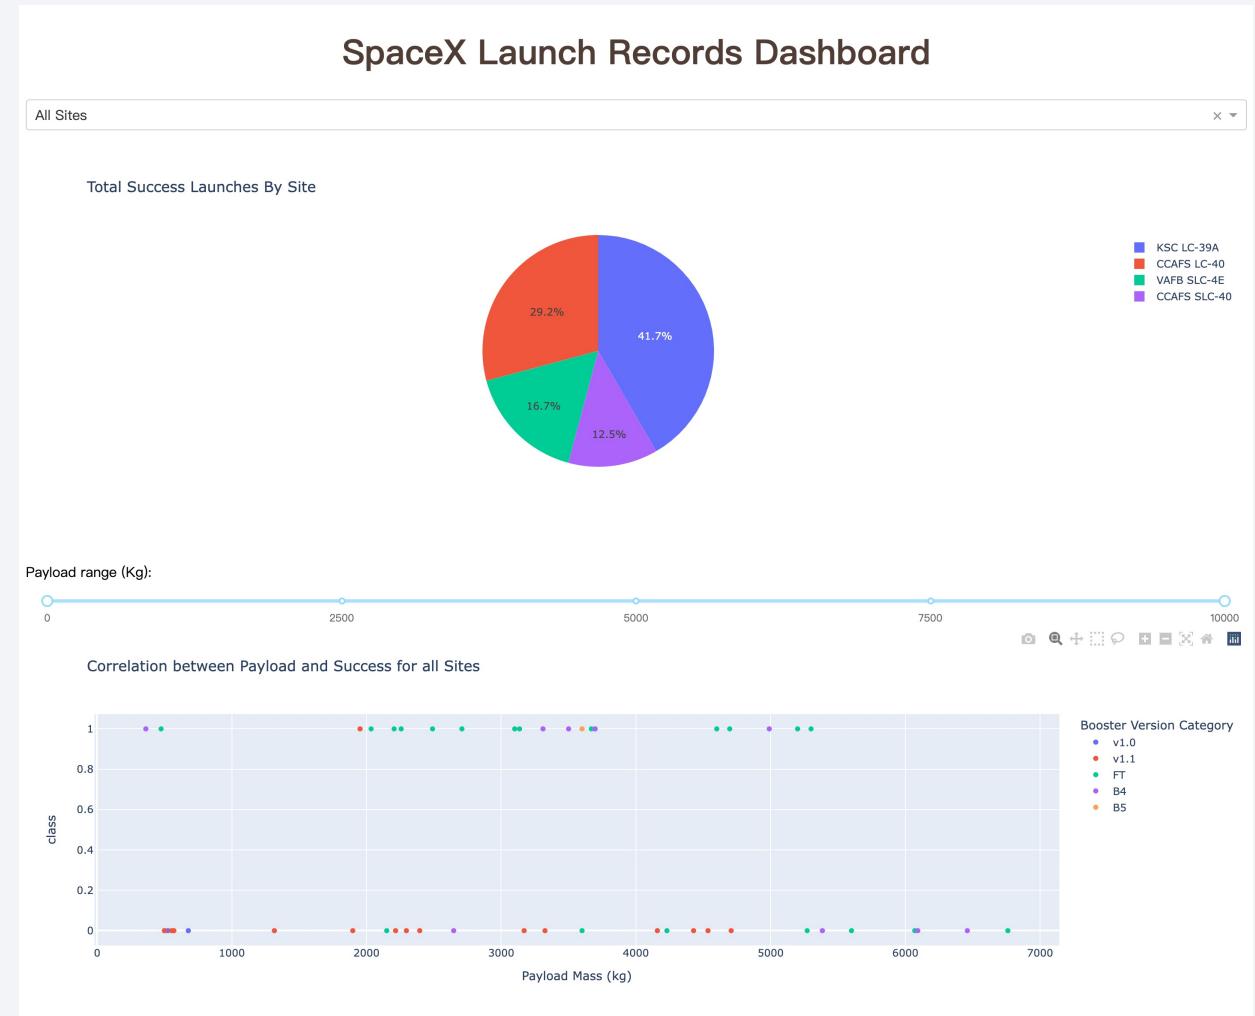


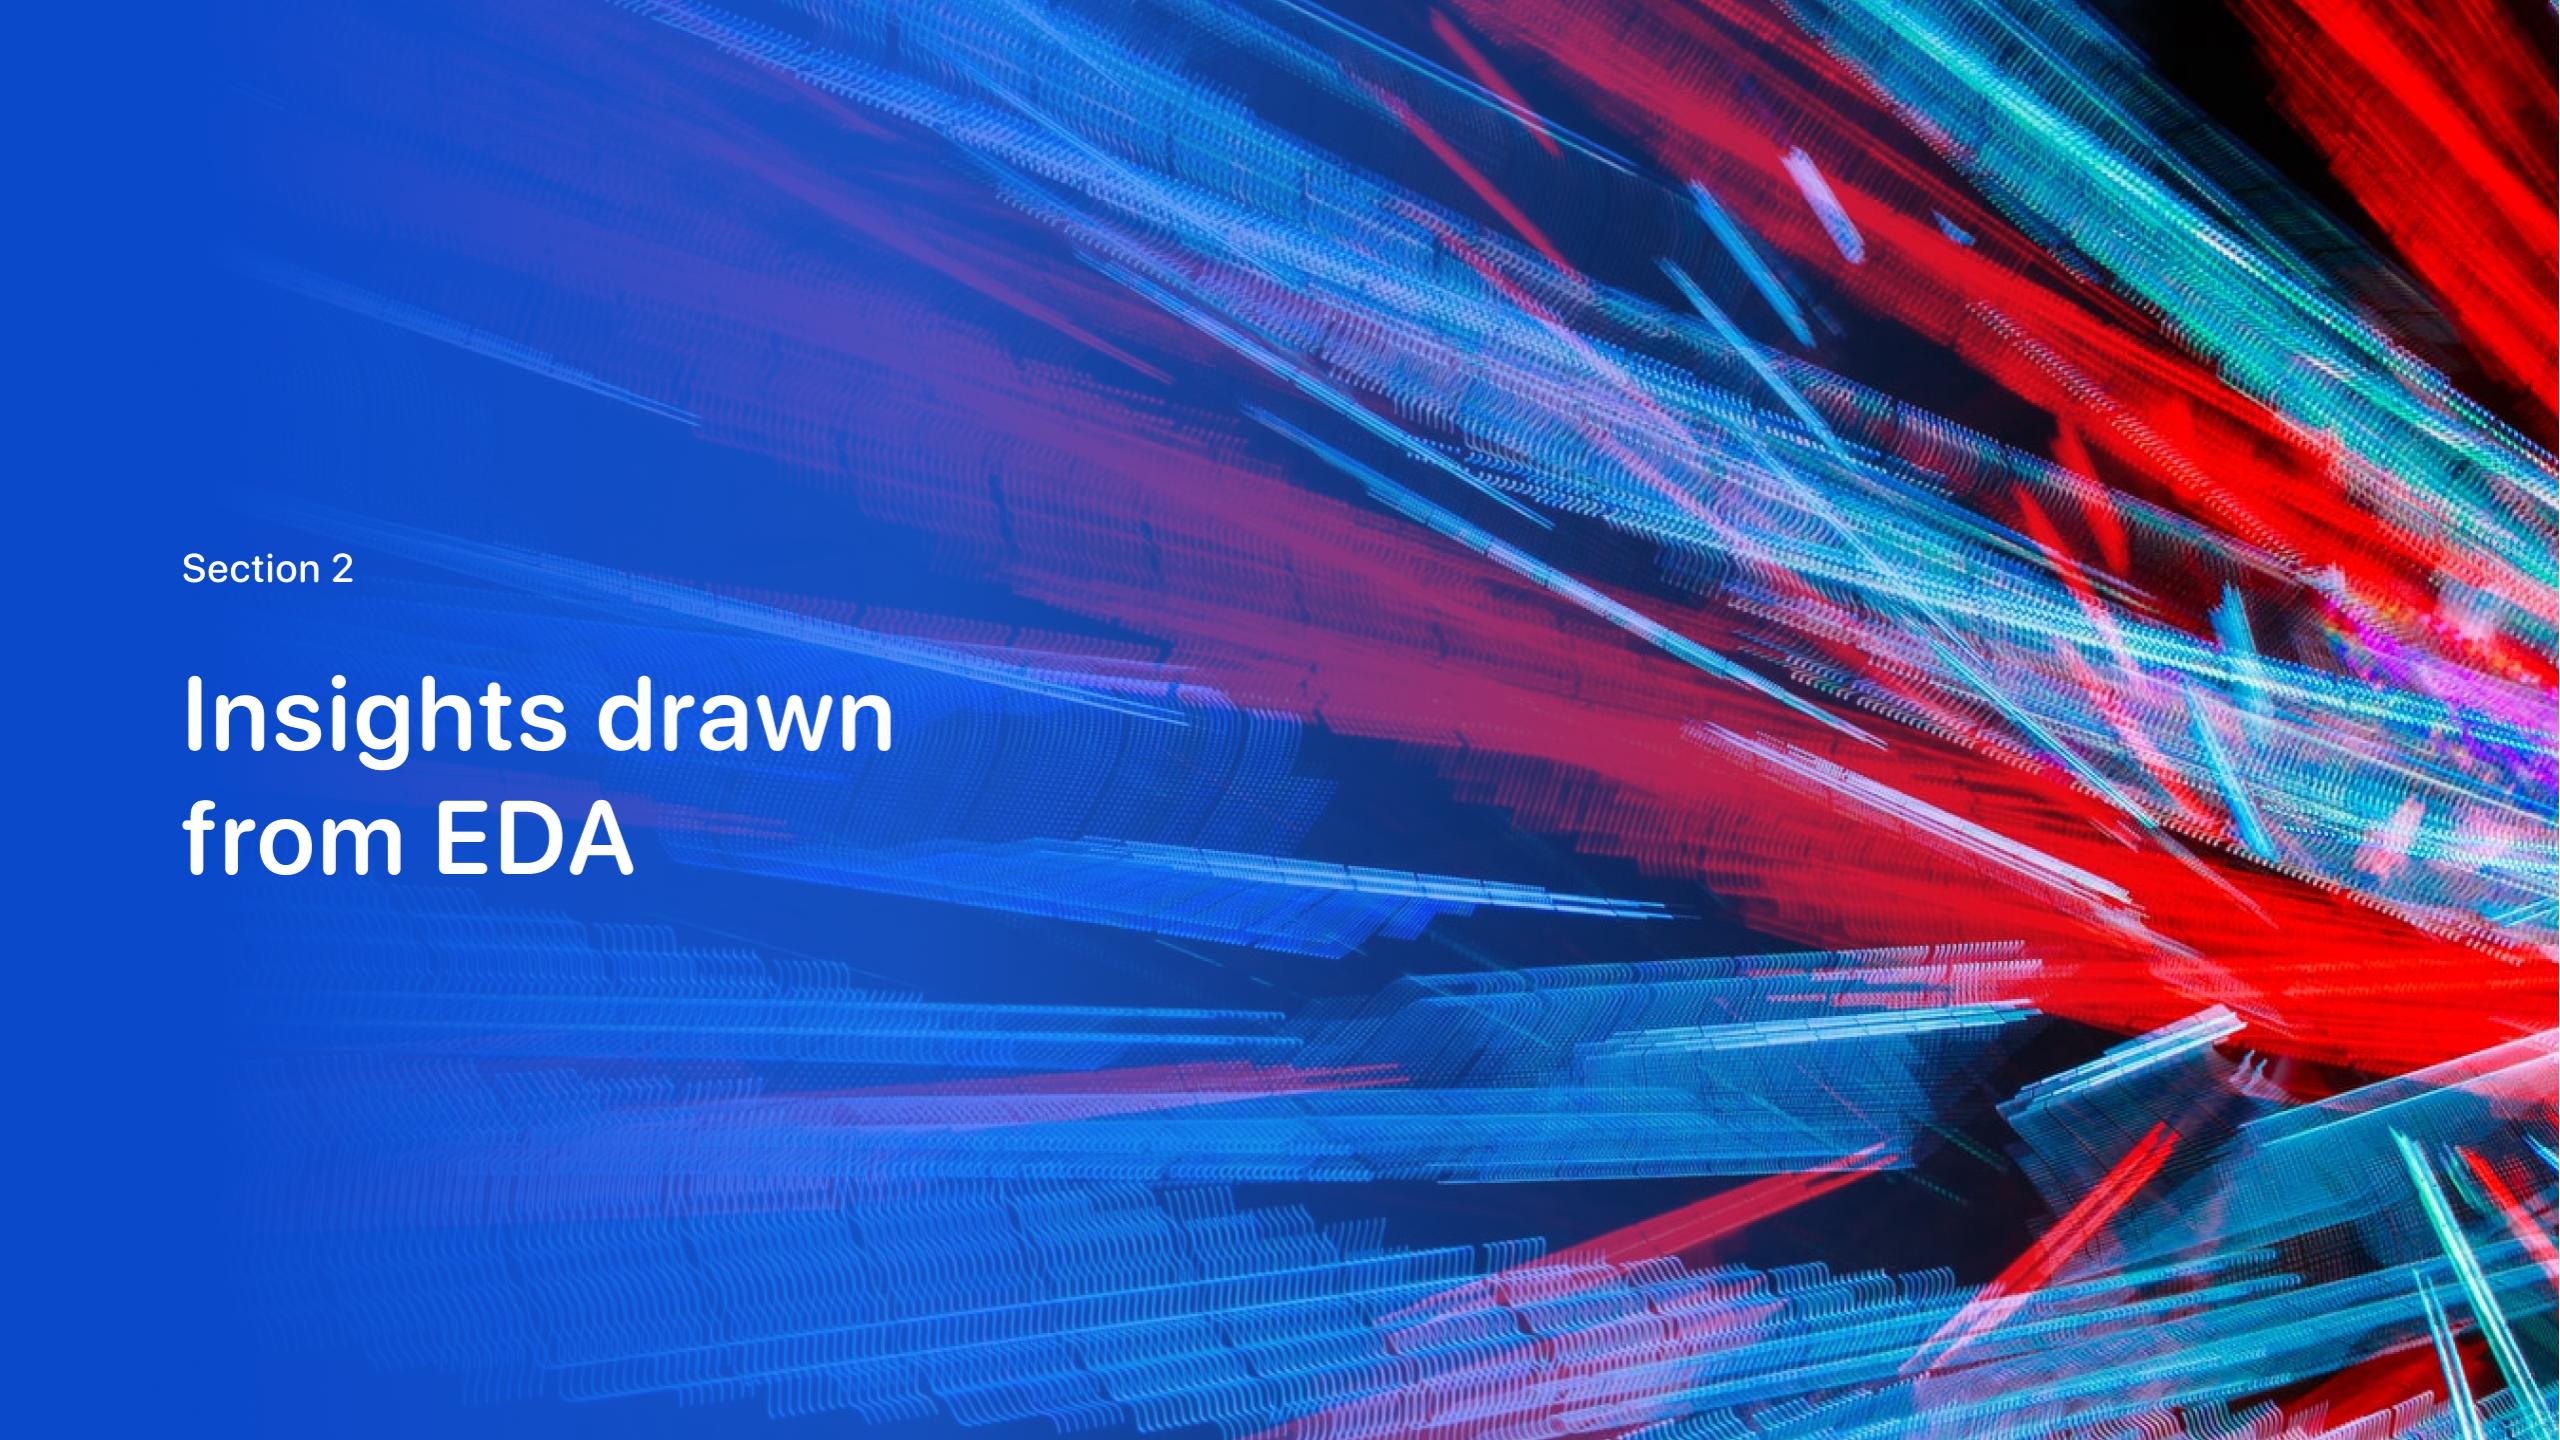
GitHub Link (Code Only):

https://github.com/sunshinekuan/IBM-Data-Science-Professional-Certificate/blob/main/Applied%20Data%20Science%20Capstone/08_Machine%20Learning%20Prediction.ipynb

Results

- The screenshot is a preview of the Dashboard with Plotly Dash.
- The results of EDA with visualization, EDA with SQL, Interactive Map with Folium, and Interactive Dashboard will be shown in the next slides.
- Comparing the accuracy of the four methods, all return the same accuracy of about 83% for test data.

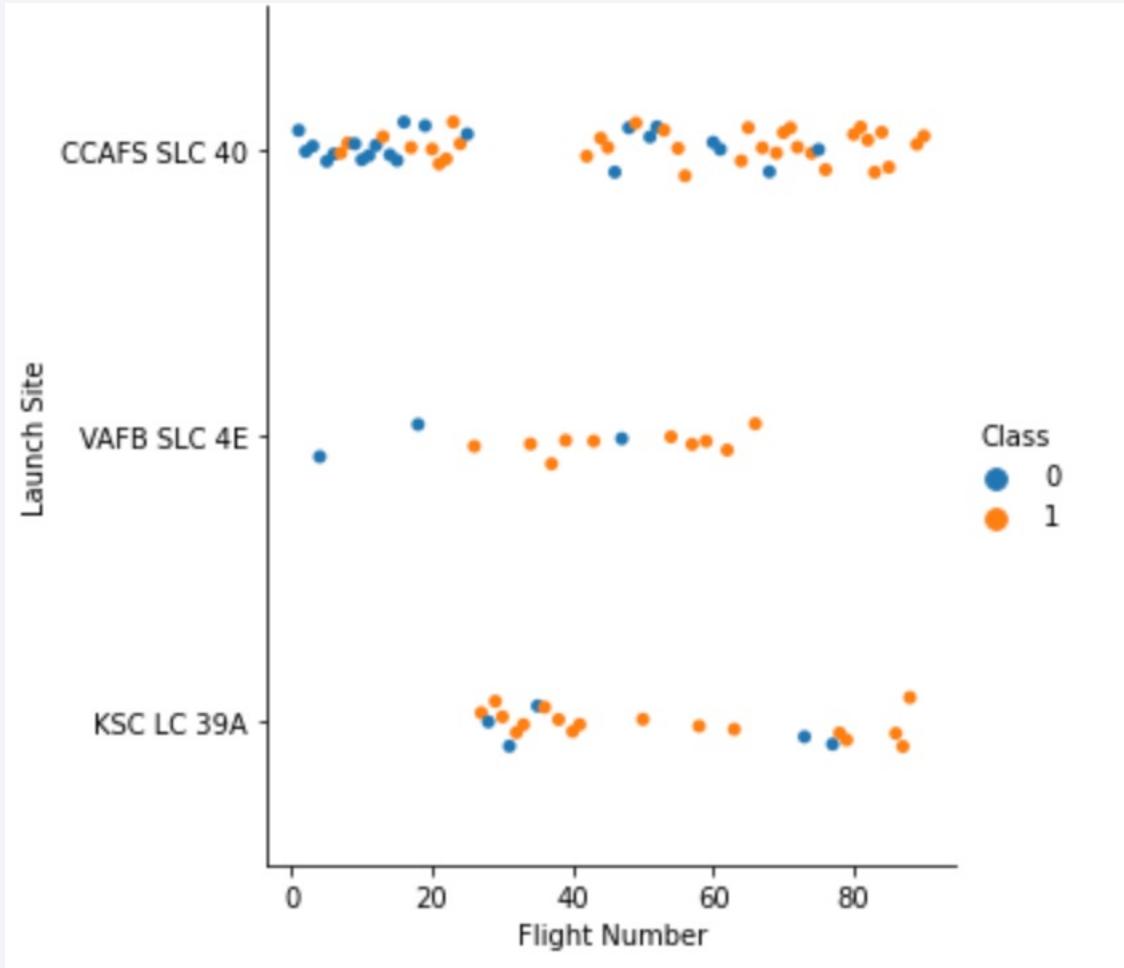


The background of the slide features a dynamic, abstract pattern of glowing particles. The particles are primarily blue and red, creating a sense of motion and depth. They are arranged in several parallel, slightly curved bands that radiate from the bottom right corner towards the top left. The intensity of the light varies, with some particles being brighter than others, which adds to the overall luminosity and three-dimensional feel of the design.

Section 2

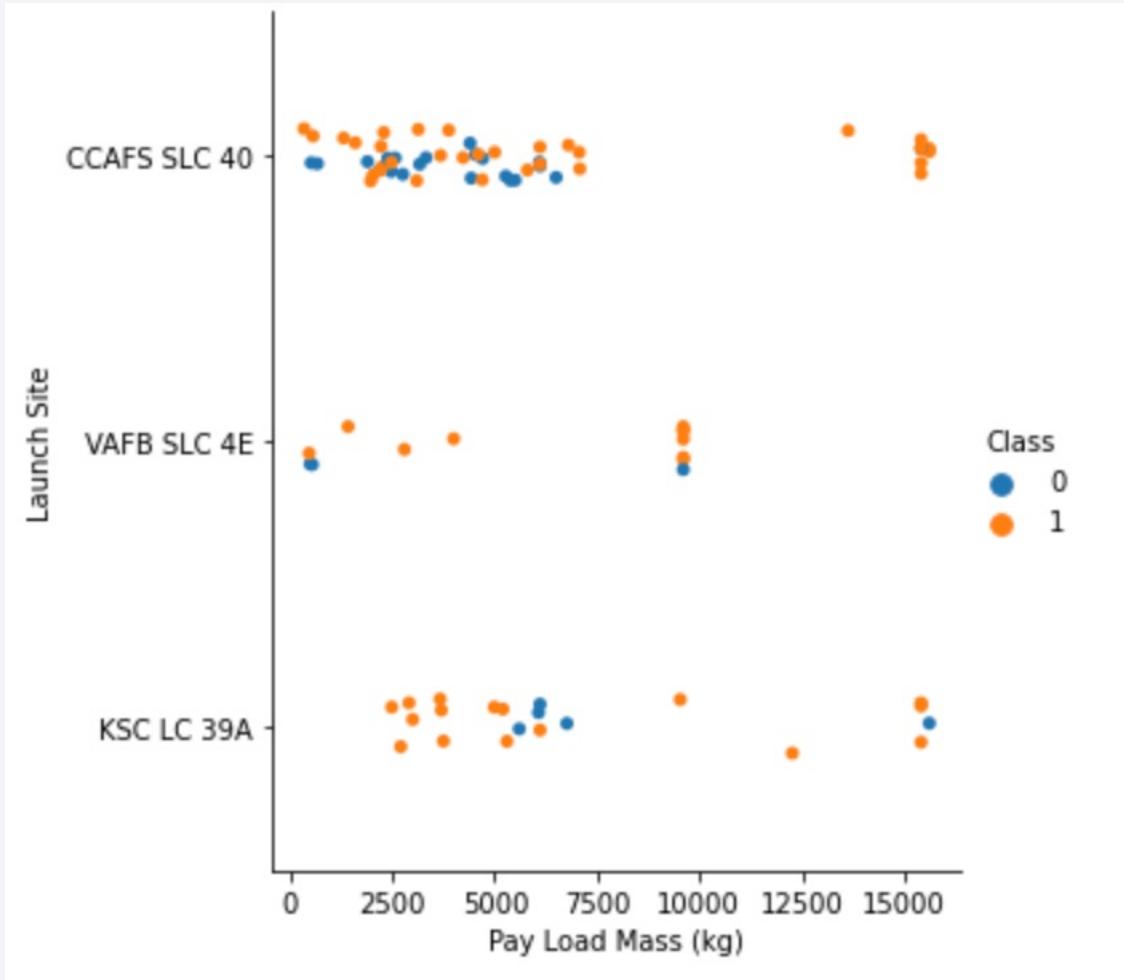
Insights drawn from EDA

Flight Number vs. Launch Site



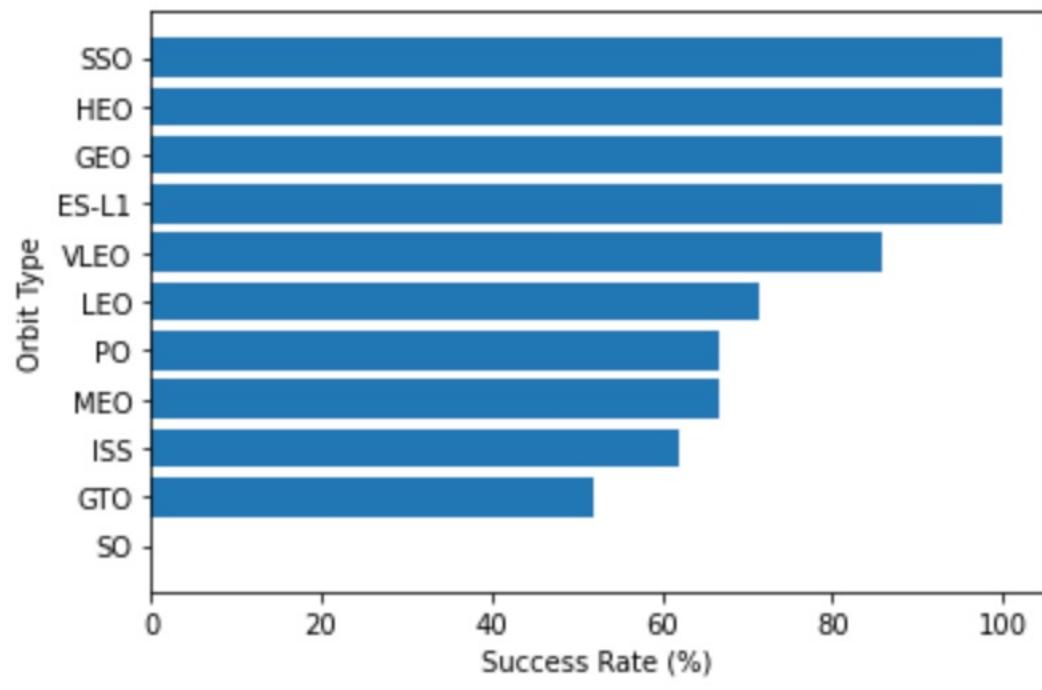
- Class 0 (blue) represents unsuccessful launch, and Class 1 (orange) represents successful launch.
- This figure shows that **the success rate increased as the number of flights increased**.
- As the success rate has increased considerably since the *20th flight*, this point seems to be a big breakthrough.

Payload vs. Launch Site



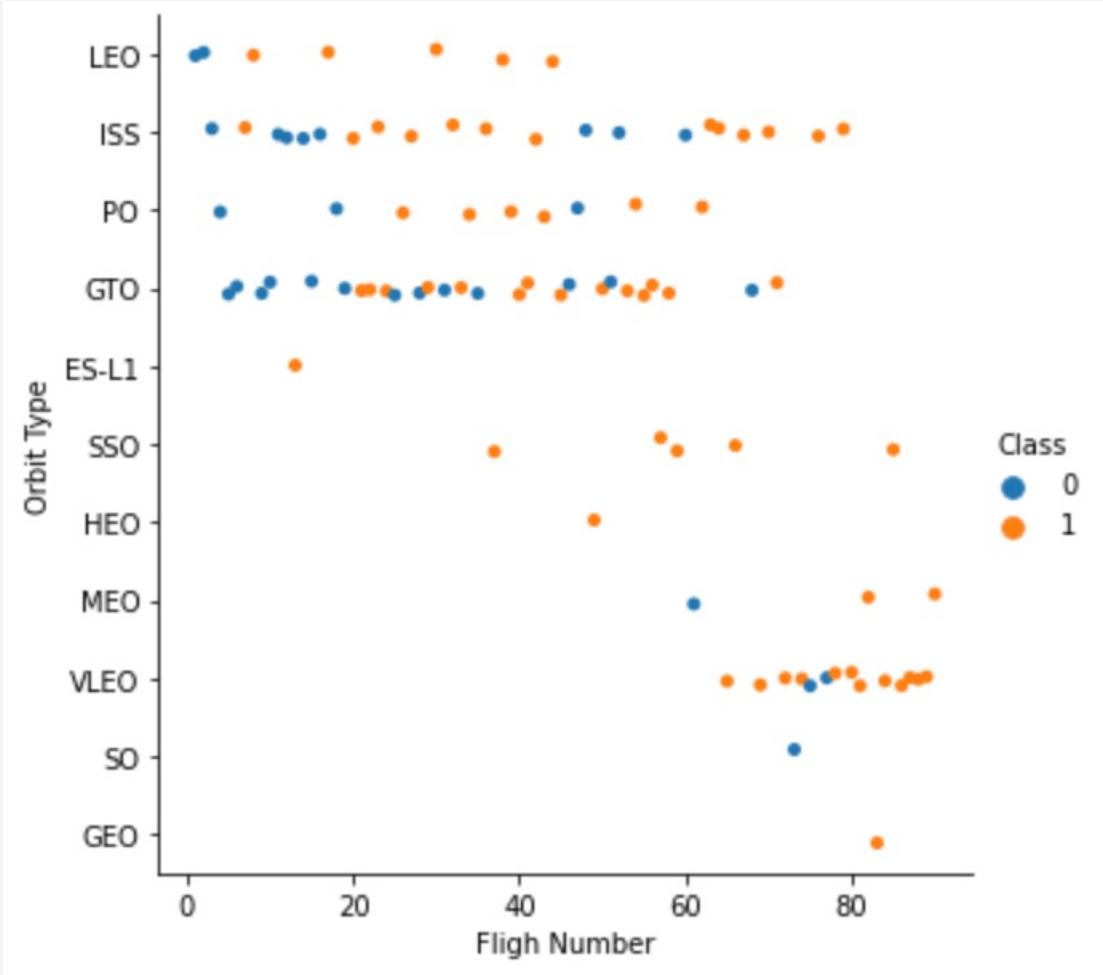
- Class 0 (blue) represents unsuccessful launch, and Class 1 (orange) represents successful launch.
- At first glance, the larger pay load mass, the higher the rocket's success rate, but it seems difficult to make decisions based on this figure because **no clear pattern can be found between successful launch and Pay Load Mass.**

Success Rate vs. Orbit Type



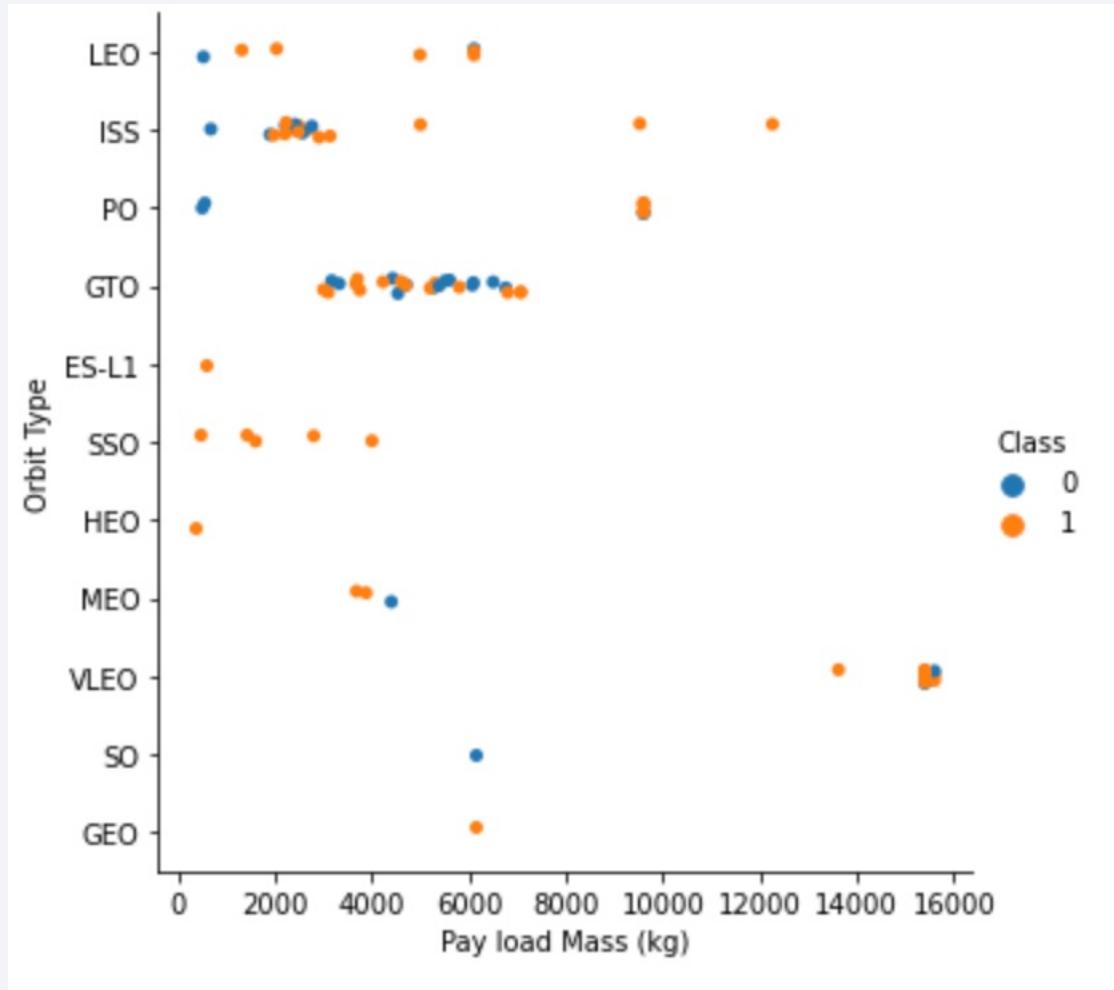
- Orbit types **SSO, HEO, GEO, and ES-L1** have **the highest success rates** (100%).
- On the other hand, the success rate of orbit type **GTO** is only 50%, and it is **the lowest** except for type SO, which recorded failure in a single attempt.

Flight Number vs. Orbit Type



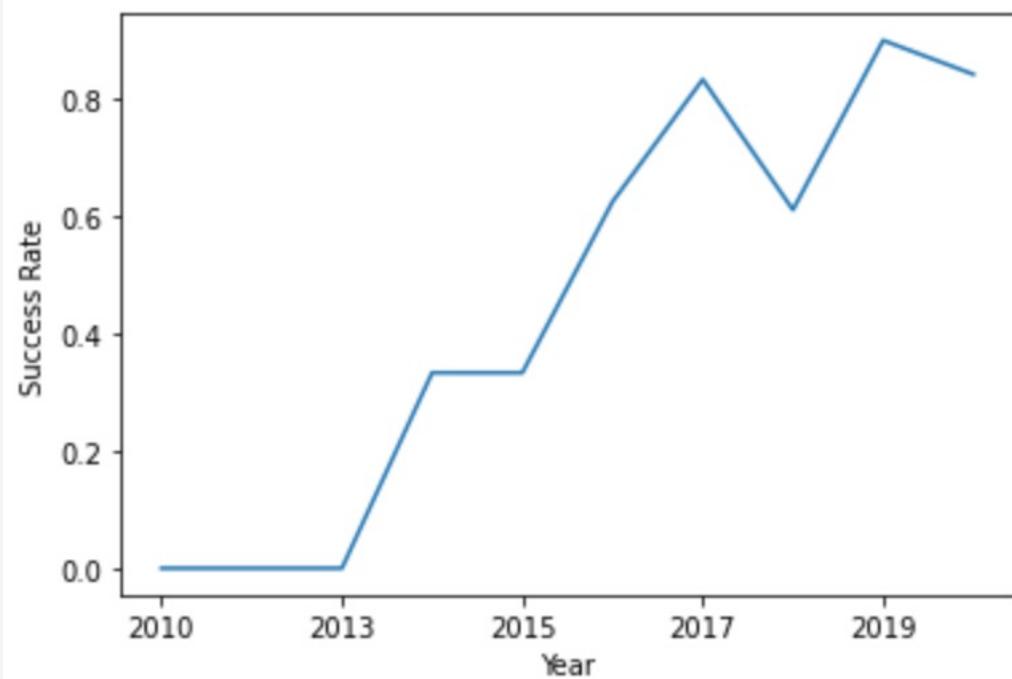
- Class 0 (blue) represents unsuccessful launch, and Class 1 (orange) represents successful launch.
- In most cases, the launch outcome seems to be correlated with the flight number.
- On the other hand, in **GTO** orbit, there seems to be **no** relationship between flight numbers and success rate.
- SpaceX starts with LEO with a moderate success rate, and it seems that VLEO, which has a high success rate, is used the most in recent launches.

Payload vs. Orbit Type



- Class 0 (blue) represents unsuccessful launch, and Class 1 (orange) represents successful launch.
- With heavy payloads the successful landing or positive landing rate are more for LEO and ISS.
- However, in the case of GTO, it is hard to distinguish between the positive landing rate and the negative landing because they are all gathering together.

Launch Success Yearly Trend



- Since 2013, the success rate has continued to **increase** until 2017.
- The rate decreased slightly in 2018.
- Recently, it has shown a success rate of about 80%.



Section 3

EDA with SQL

All Launch Site Names

Query

```
SELECT DISTINCT LAUNCH_SITE  
FROM SPACEXTBL
```

Result

launch_site

CCAFS LC-40

CCAFS SLC-40

KSC LC-39A

VAFB SLC-4E

- When the SQL **DISTINCT** clause is used in the query, only unique values are displayed in the *Launch_Site* column from the SpaceX table.
- There are four unique launch sites: CCAFS LC-40, CCAFS SLC-40, KSC LC-39A, VAFB SLC-4E

Launch Site Names Begin with 'CCA'

Query

```
SELECT * FROM SPACEXTBL  
WHERE LAUNCH_SITE LIKE 'CCA%'  
LIMIT 5
```

Result

DATE	time_utc_	booster_version	launch_site	payload	payload_mass_kg_	orbit	customer	mission_outcome	landing_outcome
2010-06-04	18:45:00	F9 v1.0 B0003	CCAFS LC-40	Dragon Spacecraft Qualification Unit	0	LEO	SpaceX	Success	Failure (parachute)
2010-12-08	15:43:00	F9 v1.0 B0004	CCAFS LC-40	Dragon demo flight C1, two CubeSats, barrel of Brouere cheese	0	LEO (ISS)	NASA (COTS) NRO	Success	Failure (parachute)
2012-05-22	07:44:00	F9 v1.0 B0005	CCAFS LC-40	Dragon demo flight C2	525	LEO (ISS)	NASA (COTS)	Success	No attempt
2012-10-08	00:35:00	F9 v1.0 B0006	CCAFS LC-40	SpaceX CRS-1	500	LEO (ISS)	NASA (CRS)	Success	No attempt
2013-03-01	15:10:00	F9 v1.0 B0007	CCAFS LC-40	SpaceX CRS-2	677	LEO (ISS)	NASA (CRS)	Success	No attempt

- Only five records of the SpaceX table were displayed using **LIMIT 5** clause in the query.
- Using the **LIKE** operator and the percent sign (%) together, the *Launch_Site* name starting with CAA could be called.

Total Payload Mass

Query

```
SELECT SUM(PAYLOAD_MASS__KG_)
AS total_payload_mass_kg
FROM SPACEXTBL
WHERE CUSTOMER = 'NASA (CRS)'
```

Result

total_payload_mass_kg
45596

- Using the **SUM()** function to calculate the sum of column **PAYLOAD_MASS__KG_**.
- In the **WHERE** clause, filter the dataset to perform calculations only if **Customer** is **NASA (CRS)**.

Average Payload Mass by F9 v1.1

Query

```
SELECT AVG(PAYLOAD_MASS__KG_)
AS avg_payload_mass_kg
FROM SPACEXTBL
WHERE BOOSTER_VERSION = 'F9 v1.1'
```

Result

avg_payload_mass_kg
2928

- Using the **AVG()** function to calculate the average value of column **PAYLOAD_MASS__KG_**.
- In the **WHERE** clause, filter the dataset to perform calculations only if **Booster_version** is **F9 v1.1**.

First Successful Ground Landing Date

Query

```
SELECT MIN(DATE)
AS first_successful_landing_date
FROM SPACEXTBL
WHERE LANDING_OUTCOME =
'Success (ground pad)'
```

- Using the **MIN()** function to find out the earliest date in the column **DATE**.
- In the **WHERE** clause, filter the dataset to perform a search only if **Landing_outcome** is **Success (ground pad)**.

Result

first_successful_landing_date

2015-12-22

Successful Drone Ship Landing with Payload between 4000 and 6000

Query

```
SELECT BOOSTER_VERSION
FROM SPACEXTBL
WHERE LANDING_OUTCOME =
'Success (drone ship)'
AND (PAYLOAD_MASS_KG_
BETWEEN 4000 AND 6000)
```

- In the **WHERE** clause, filter the dataset to perform a search if **Landing_outcome** is **Success (drone ship)**.
 - Using the **AND** operator to display a record if additional condition **PAYLOAD_MASS_KG_** is between **4000** and **6000**.

Result

	booster_version
	F9 FT B1022
	F9 FT B1026
	F9 FT B1021.2
	F9 FT B1031.2

Total Number of Successful and Failure Mission Outcomes

Query

```
SELECT MISSION_OUTCOME, COUNT(*)  
AS total_number  
FROM SPACEXTBL  
GROUP BY MISSION_OUTCOME
```

Result

mission_outcome	total_number
Failure (in flight)	1
Success	99
Success (payload status unclear)	1

- Using the **COUNT()** function to calculate the total number of columns.
- Using the **GROUP BY** statement, groups rows that have the same values into summary rows to find the total number in each **Mission_outcome**.
- According to the result, SpaceX seems to have **successfully completed nearly 99% of its missions**.

Boosters Carried Maximum Payload

Query

```
SELECT DISTINCT BOOSTER_VERSION,
PAYLOAD_MASS__KG_
FROM SPACEXTBL
WHERE PAYLOAD_MASS__KG_ =
( SELECT MAX(PAYLOAD_MASS__KG_)
FROM SPACEXTBL);
```

Result

booster_version	payload_mass_kg_
F9 B5 B1048.4	15600
F9 B5 B1048.5	15600
F9 B5 B1049.4	15600
F9 B5 B1049.5	15600
F9 B5 B1049.7	15600
F9 B5 B1051.3	15600
F9 B5 B1051.4	15600
F9 B5 B1051.6	15600
F9 B5 B1056.4	15600
F9 B5 B1058.3	15600
F9 B5 B1060.2	15600
F9 B5 B1060.3	15600

- Using a subquery, first, find the maximum value of the payload by using MAX() function, and second, filter the dataset to perform a search if PAYLOAD_MASS__KG_ is the maximum value of the payload.
- According to the result, version F9 B5 B10xx.x boosters could carry the maximum payload.

2015 Launch Records

Query

```
SELECT LANDING__OUTCOME,  
BOOSTER_VERSION, LAUNCH_SITE  
FROM SPACEXTBL  
WHERE LANDING__OUTCOME = 'Failure  
(drone ship) AND YEAR(DATE) = '2015'
```

Result

landing__outcome	booster_version	launch_site
Failure (drone ship)	F9 v1.1 B1012	CCAFS LC-40
Failure (drone ship)	F9 v1.1 B1015	CCAFS LC-40

- In the **WHERE** clause, filter the dataset to perform a search if **Landing__outcome** is **Failure (drone ship)**.
- Using the **AND** operator to display a record if additional condition **YEAR** is 2015.
- In 2015, there were two landing failures on drone ships.

Rank Landing Outcomes Between 2010-06-04 and 2017-03-20

Query

```
SELECT LANDING_OUTCOME,  
COUNT(LANDING_OUTCOME)  
AS total_number  
FROM SPACEXTBL  
WHERE DATE BETWEEN '2010-06-04'  
AND '2017-03-20'  
GROUP BY LANDING_OUTCOME  
ORDER BY total_number DESC
```

Result

landing_outcome	total_number
No attempt	10
Failure (drone ship)	5
Success (drone ship)	5
Controlled (ocean)	3
Success (ground pad)	3
Failure (parachute)	2
Uncontrolled (ocean)	2
Precluded (drone ship)	1

- In the **WHERE** clause, filter the dataset to perform a search if the date is between **2010-06-04** and **2017-03-20**.
- Using the **ORDER BY** keyword to sort the records by total number of landing and using **DESC** keyword to sort the records in descending order.
- According to the results, the number of successes and failures between 2010-06-04 and 2017-03-20 was similar.

The background of the slide is a photograph taken from space at night. It shows the curvature of the Earth against a dark blue-black void of space. City lights are visible as numerous small white and yellow dots, primarily concentrated in the lower right quadrant where the United States appears. In the upper left quadrant, the green and yellow glow of the Aurora Borealis (Northern Lights) is visible in the upper atmosphere.

Section 4

Launch Sites Proximities Analysis

<Folium Map Screenshot 1>

- Replace <Folium map screenshot 1> title with an appropriate title
- Explore the generated folium map and make a proper screenshot to include all launch sites' location markers on a global map
- Explain the important elements and findings on the screenshot

<Folium Map Screenshot 2>

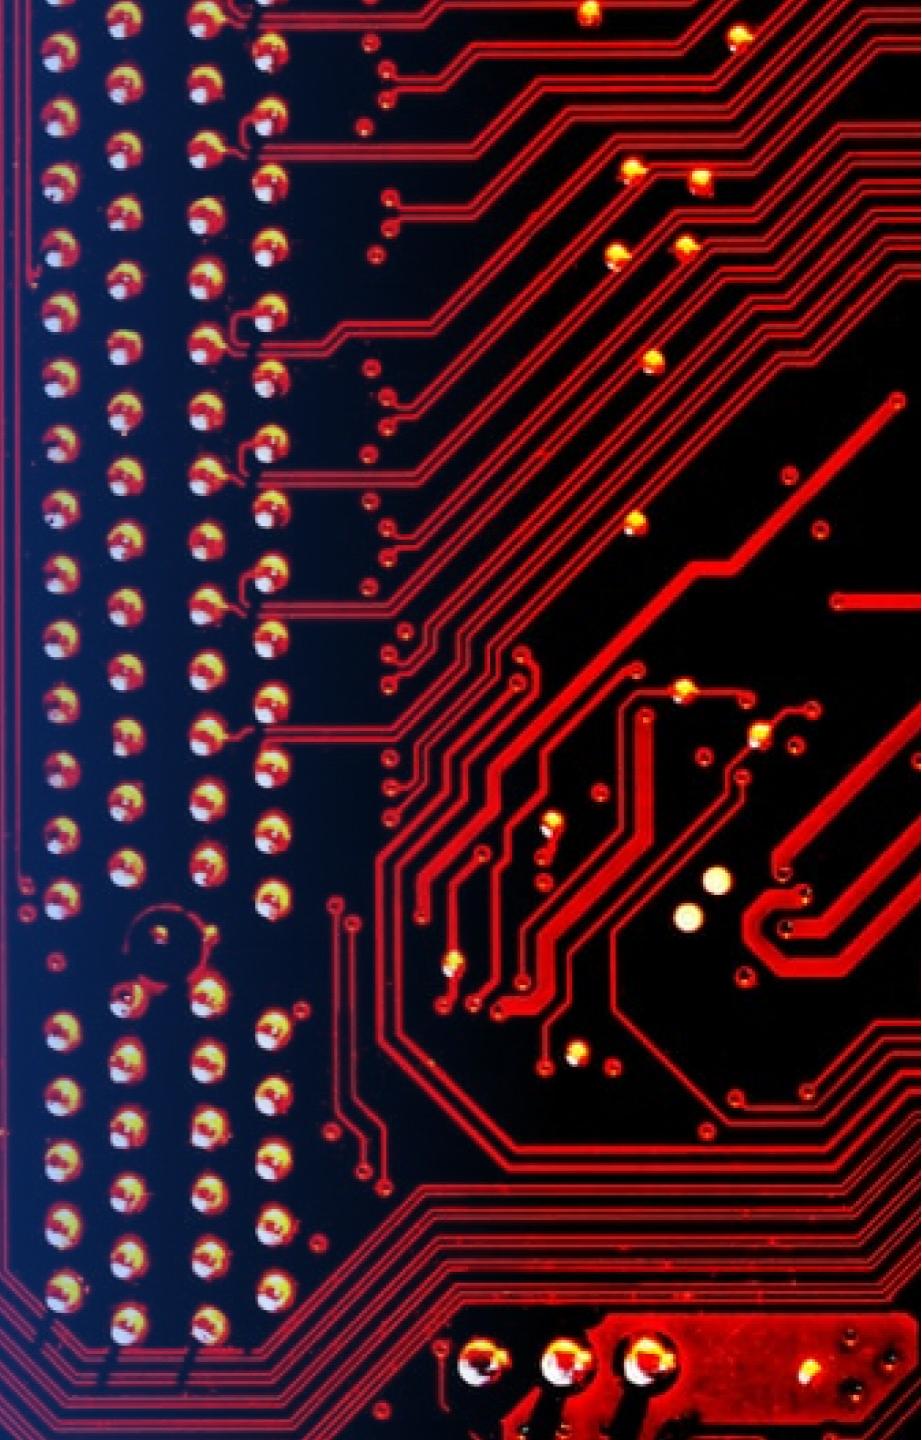
- Replace <Folium map screenshot 2> title with an appropriate title
- Explore the folium map and make a proper screenshot to show the color-labeled launch outcomes on the map
- Explain the important elements and findings on the screenshot

<Folium Map Screenshot 3>

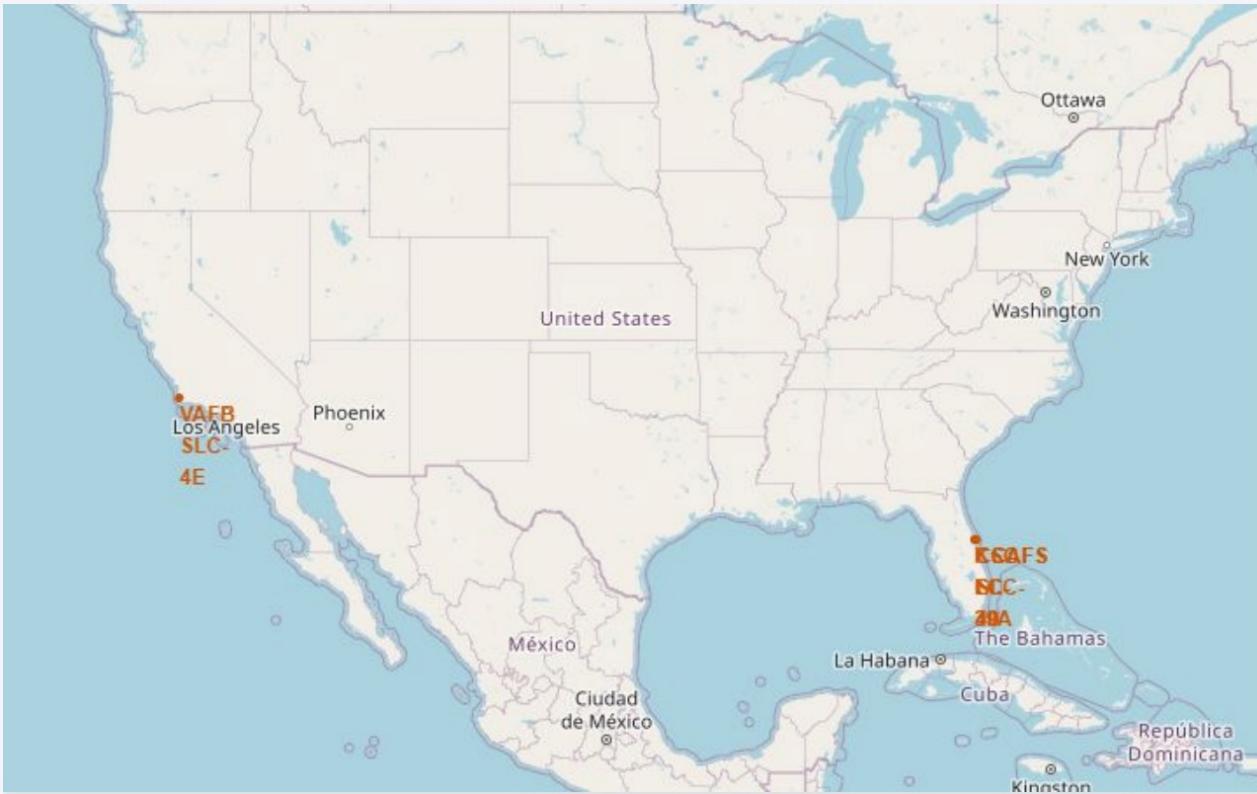
- Replace <Folium map screenshot 3> title with an appropriate title
- Explore the generated folium map and show the screenshot of a selected launch site to its proximities such as railway, highway, coastline, with distance calculated and displayed
- Explain the important elements and findings on the screenshot

Section 5

Build a Dashboard with Plotly Dash

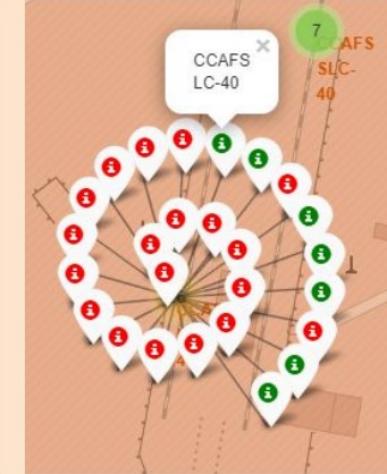
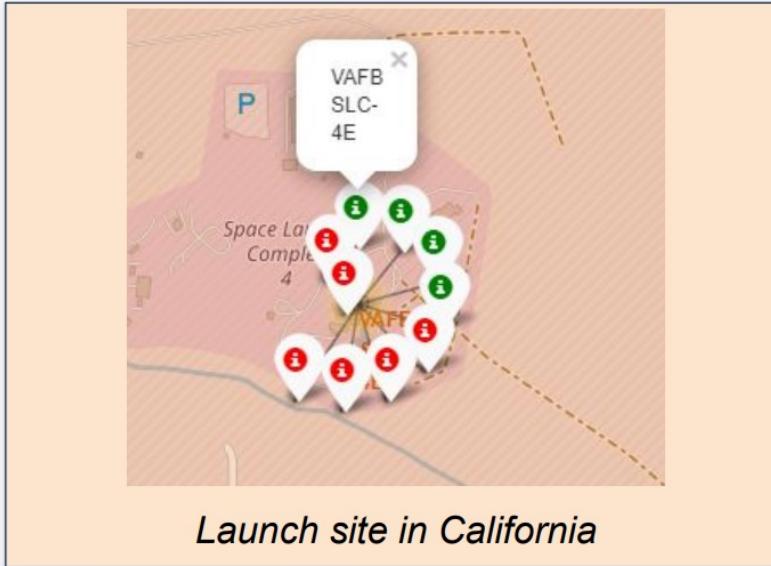


<Dashboard Screenshot 1>



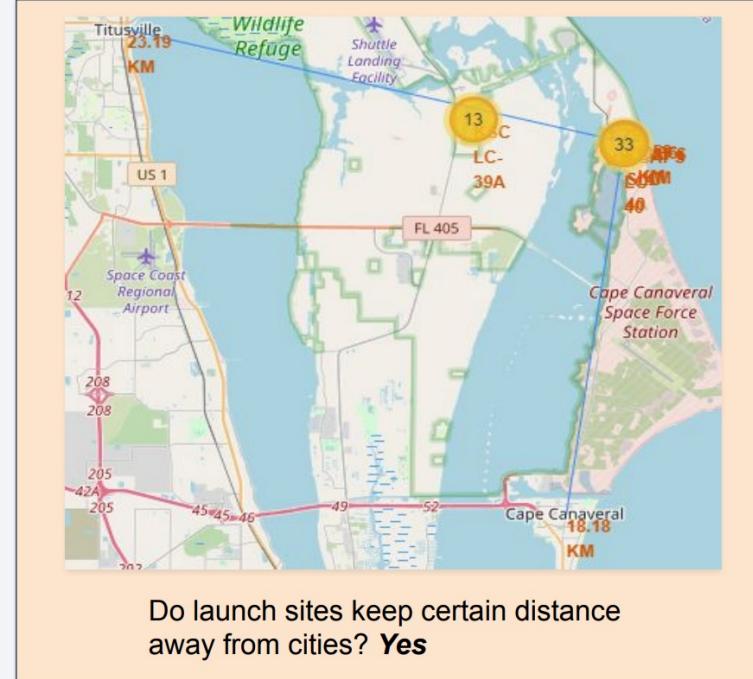
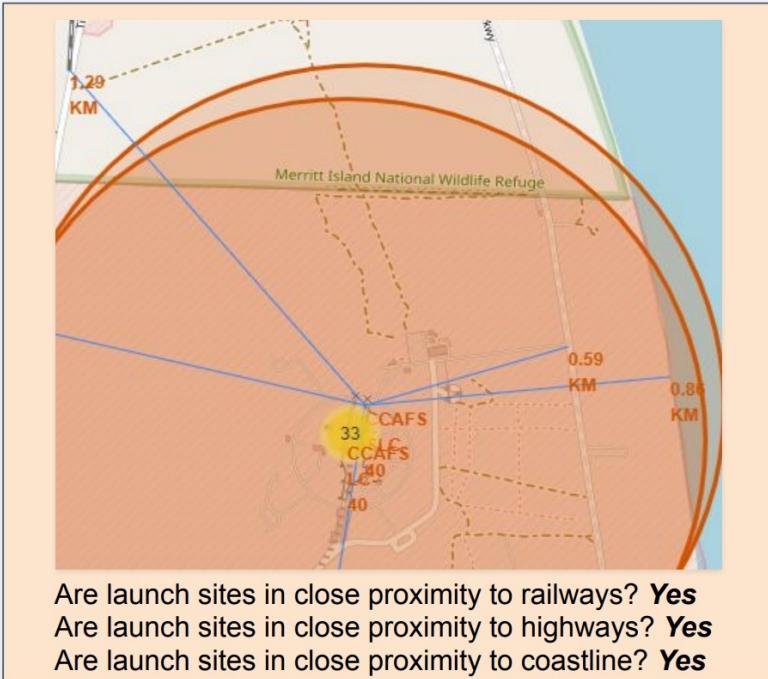
- The left map shows all *SpaceX* launch sites, and the right map also shows that all launch sites are in the United States.
- As can be seen on the map, all launch sites are near the coast.

<Dashboard Screenshot 2>



- By clicking on the marker clusters, successful landing (**green**) or failed landing (**red**) are displayed.

<Dashboard Screenshot 3>

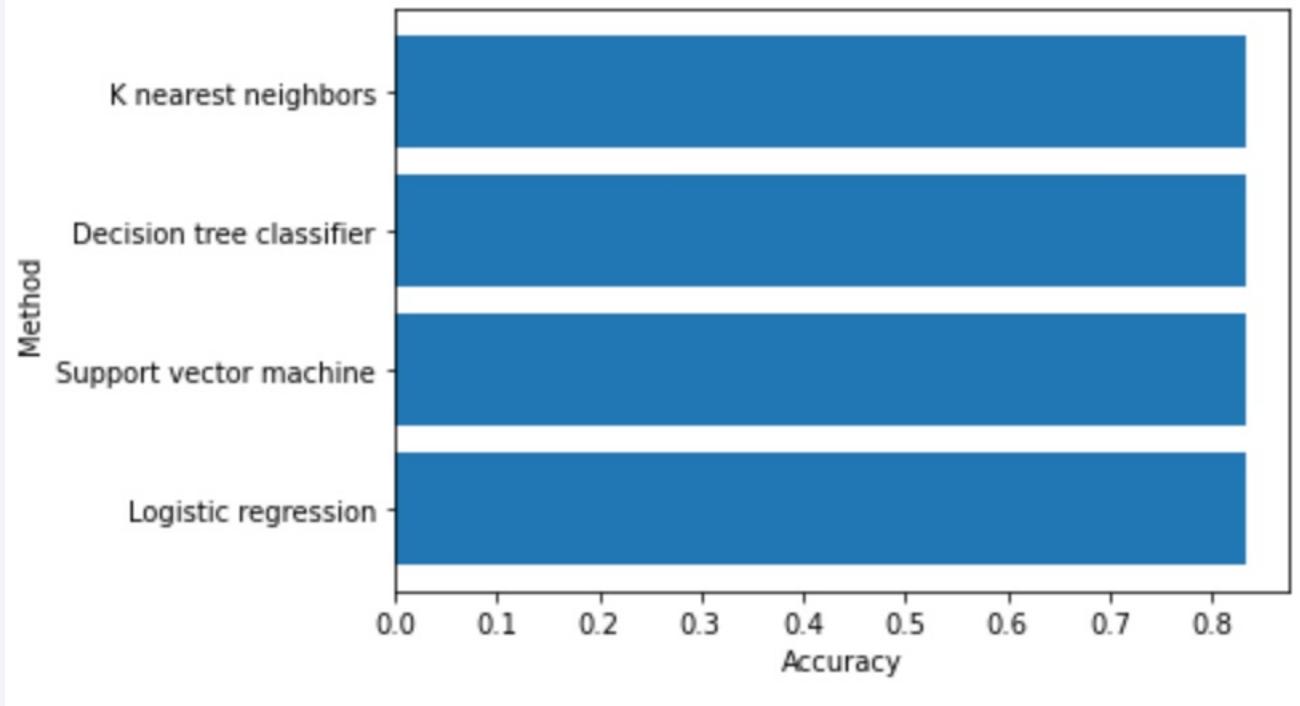


- It can be found that the launch site is **close to railways and highways** for transportation of equipment or personnel and is also **close to coastline** and relatively **far from the cities** so that launch failure does not pose a threat.

Section 6

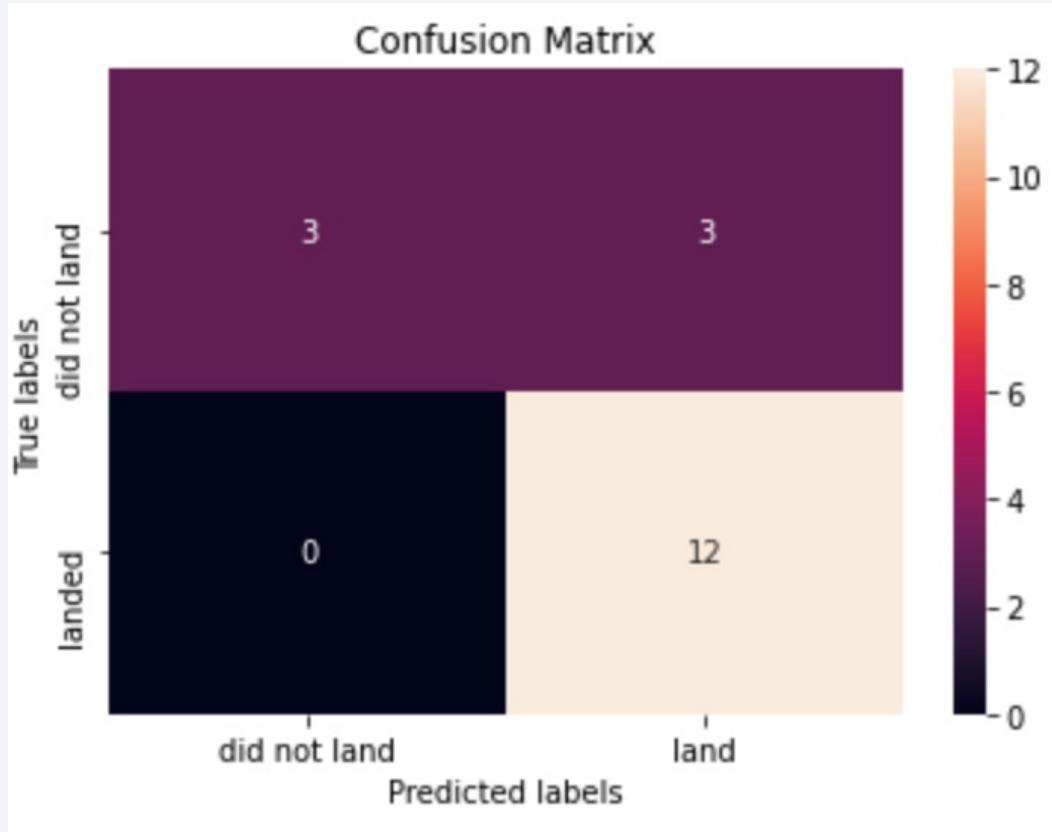
Predictive Analysis (Classification)

Classification Accuracy



- In the test set, **the accuracy of all models** was virtually the **same** at **83.33%**.
- It should be noted that the test size was small at 18.
- Therefore, more data is needed to determine the optimal model.

Confusion Matrix



- The confusion matrix is the same for all models because all models performed the same for the test set.
- The models predicted 12 successful landings when the true label was successful and 3 failed landings when the true label was failure. But there were also 3 predictions that said successful landings when the true label was failure (*false positive*).
- Overall, **these models predict successful landings**.

Conclusions

- As the number of flights increased, the success rate increased, and recently it has exceeded 80%.
- Orbital types SSO, HEO, GEO, and ES-L1 have the highest success rate (100%).
- The launch site is close to railways, highways, and coastline, but far from cities.
- KSLC-39A has the highest number of launch successes and the highest success rate among all sites.
- The launch success rate of low weighted payloads is higher than that of heavy weighted payloads.
- In this dataset, all models have the same accuracy (83.33%), but it seems that more data is needed to determine the optimal model due to the small data size.

Appendix

- GitHub Link
 - <https://github.com/sunshinekuan/IBM-Data-Science-Professional-Certificate/tree/main/Applied%20Data%20Science%20Capstone>
- Course Link – Coursera, IBM - Applied Data Science Capstone
 - <https://www.coursera.org/learn/applied-data-science-capstone/home/welcome>

Thank you!

